

# CO324: TCP clients and servers

Ziyan Maraikar

February 5, 2015

# Lecture Outline

1 TCP Clients

2 TCP Servers

3 Framing

TCP is a *stream protocol* in which packet boundaries are invisible to the application. Data is received and transmitted as a sequence of bytes.

It provides

- 1 Reliability

TCP is a *stream protocol* in which packet boundaries are invisible to the application. Data is received and transmitted as a sequence of bytes.

It provides

- 1 Reliability
- 2 Ordering

TCP is a *stream protocol* in which packet boundaries are invisible to the application. Data is received and transmitted as a sequence of bytes.

It provides

- ➊ Reliability
- ➋ Ordering
- ➌ Flow control

# TCP clients

The Java `Socket` class represents a TCP socket.

```
Socket socket = new Socket();
```

# TCP clients

The Java `Socket` class represents a TCP socket.

```
Socket socket = new Socket();
```

```
InputStream sin = socket.getInputStream();
```

```
OutputStream sout = socket.getOutputStream();
```

Binary data is read and written to the socket via the associated I/O streams.

# Sending and receiving text

```
try {  
    Socket socket = new Socket(address, PORT);  
  
    BufferedReader sin = new BufferedReader (  
        new InputStreamReader(socket.getInputStream() ));  
  
    BufferedWriter sout = new BufferedWriter (  
        new OutputStreamWriter(socket.getOutputStream() ));  
} catch(IOException e) {  
    e.printStackTrace();  
}
```

Note that application messages must be properly *framed* e.g. using a delimiter like `\n`.



# Sending and receiving text

```
try {  
    Socket socket = new Socket(address, PORT);  
  
    BufferedReader sin = new BufferedReader (  
        new InputStreamReader(socket.getInputStream() ));  
  
    BufferedWriter sout = new BufferedWriter (  
        new OutputStreamWriter(socket.getOutputStream() ));  
} catch(IOException e) {  
    e.printStackTrace();  
}
```

Note that application messages must be properly *framed* e.g. using a delimiter like `\n`.

What do the `BufferedReader` and `BufferedWriter` classes do?

# Lecture Outline

1 TCP Clients

2 TCP Servers

3 Framing

# TCP Servers

Java uses a separate `ServerSocket` class to bind to a port and accept connections from clients.

```
ServerSocket ss = new ServerSocket(PORT);  
Socket socket = ss.accept();
```

`accept` returns a new socket connected to the client.

# Server example

```
ServerSocket ss = new ServerSocket(PORT);
Socket socket = ss.accept();

while (true)
try {
    BufferedReader sin = new BufferedReader (
        new InputStreamReader(socket.getInputStream() ));

    BufferedWriter sout = new BufferedWriter (
        new OutputStreamWriter(socket.getOutputStream() ));

} catch(IOException e) {
    e.printStackTrace();
}
```

# Server example

```
ServerSocket ss = new ServerSocket(PORT);
Socket socket = ss.accept();

while (true)
try {
    BufferedReader sin = new BufferedReader (
        new InputStreamReader(socket.getInputStream() ));

    BufferedWriter sout = new BufferedWriter (
        new OutputStreamWriter(socket.getOutputStream() ));
} catch(IOException e) {
    e.printStackTrace();
}
```

What happens if multiple clients try to connect at once?

# Lecture Outline

1 TCP Clients

2 TCP Servers

**3 Framing**

# Messages on streams

Suppose a client sends two consecutive messages, and the server does a **read**. Will it get

- ① both messages at once?
- ② the first message only?
- ③ part of the first message?

# Messages on streams

Suppose a client sends two consecutive messages, and the server does a **read**. Will it get

- ① both messages at once?
- ② the first message only?
- ③ part of the first message?

It depends on network conditions and the TCP/IP stacks!



TCP can only send to and receive from a byte stream, but application protocols are built with discrete messages.

# Framing

TCP can only send to and receive from a byte stream, but application protocols are built with discrete messages.

We must define a method of *framing* application messages, so that message boundaries are unambiguous.

Method used depends on the kind of data

- ★ Text
- ★ Binary

# Text protocols

Most application protocols on the Internet are textual.

- ★ Human readable — so easy to debug.

# Text protocols

Most application protocols on the Internet are textual.

- ★ Human readable — so easy to debug.
- ★ Historically, the most applications were textual e.g., Telnet, Email

# Text protocols

Most application protocols on the Internet are textual.

- ★ Human readable — so easy to debug.
- ★ Historically, the most applications were textual e.g., Telnet, Email

# Text protocols

Most application protocols on the Internet are textual.

- ★ Human readable — so easy to debug.
- ★ Historically, the most applications were textual e.g., Telnet, Email

Disadvantages:

- ★ Vulnerable to security attacks like buffer overflows.

# Text protocols

Most application protocols on the Internet are textual.

- ★ Human readable — so easy to debug.
- ★ Historically, the most applications were textual e.g., Telnet, Email

Disadvantages:

- ★ Vulnerable to security attacks like buffer overflows.
- ★ Bandwidth and CPU inefficient.

## Example: SMTP

```
S: 220 smtp.server.com Simple Mail Transfer Service Ready
C: HELO client.example.com
S: 250 Hello client.example.com
C: MAIL FROM:<jane@yahoo.com>
S: 250 OK
C: RCPT TO:<john@gmail.com>
S: 250 OK
C: DATA
S: 354 Send message content; end with <CRLF>.<CRLF>
C: <The message data (body text, subject, e-mail header,
C: .
S: 250 OK, message accepted for delivery: queued as 123456789
C: QUIT
S: 221 Bye
```



## Example: SMTP

```
S: 220 smtp.server.com Simple Mail Transfer Service Ready
C: HELO client.example.com
S: 250 Hello client.example.com
C: MAIL FROM:<jane@yahoo.com>
S: 250 OK
C: RCPT TO:<john@gmail.com>
S: 250 OK
C: DATA
S: 354 Send message content; end with <CRLF>.<CRLF>
C: <The message data (body text, subject, e-mail header,
C: .
S: 250 OK, message accepted for delivery: queued as 1234567890
C: QUIT
S: 221 Bye
```

How are non-text mail attachments handled?

# Delimiters

We can *delimit* text protocol messages using a special character.

The usual delimiter used in Internet protocols are the line termination characters **CR**, **LF** or **CRLF**.

# Delimiters

We can *delimit* text protocol messages using a special character.

The usual delimiter used in Internet protocols are the line termination characters **CR**, **LF** or **CRLF**.

```
Socket socket = new Socket(address, PORT);

BufferedReader sin = new BufferedReader (
    new InputStreamReader(socket.getInputStream() ));

BufferedWriter sout = new BufferedWriter (
    new OutputStreamWriter(socket.getOutputStream() ));

sout.write("hello world\n");
sin.readLine();
```

`BufferedReader.readLine` splits the apart newline delimited messages in a stream.

# Binary protocols

Binary protocols support transmission of arbitrary data. Usually contains a fixed-format *header* that describes the *payload*.

# Binary protocols

Binary protocols support transmission of arbitrary data. Usually contains a fixed-format *header* that describes the *payload*.

- ★ Suited to describing structured data.

# Binary protocols

Binary protocols support transmission of arbitrary data. Usually contains a fixed-format *header* that describes the *payload*.

- ★ Suited to describing structured data.
- ★ Easier to *parse* — metadata received before payload.

# Binary protocols

Binary protocols support transmission of arbitrary data. Usually contains a fixed-format *header* that describes the *payload*.

- ★ Suited to describing structured data.
- ★ Easier to *parse* — metadata received before payload.
- ★ Efficient use of bandwidth.

# Binary protocols

Binary protocols support transmission of arbitrary data. Usually contains a fixed-format *header* that describes the *payload*.

- ★ Suited to describing structured data.
- ★ Easier to *parse* — metadata received before payload.
- ★ Efficient use of bandwidth.



# Binary protocols

Binary protocols support transmission of arbitrary data. Usually contains a fixed-format *header* that describes the *payload*.

- ★ Suited to describing structured data.
- ★ Easier to *parse* — metadata received before payload.
- ★ Efficient use of bandwidth.

Example: Basic encoding rules for ASN.1, an OSI standard used in protocols such as LDAP.

|      |        |       |                |
|------|--------|-------|----------------|
| Type | Length | Value | End-of-content |
|------|--------|-------|----------------|

# Example

```
Socket socket = new Socket(address, PORT);

DataInputStream sin = new DataInputStream(
    socket.getInputStream() );
DataOutputStream sout = new DataOutputStream(
    socket.getOutputStream())

byte[] data = new byte [100000];
data[data.length-1] = (byte) 127;
sout.writeInt(data.length);
sout.write(data);
```